

Models

→ Saving models (Pytorch, sklearn...)

MLflow - Projects

4-11-2025

1 → python 1_tracking.py

2 → mlflow run . -e tracking --env-manager local

↳ local (can give github repo as well)

- Command 1 ran, but command 2 didn't. Since mlflow requires some more data which we didn't give. in command.
- We actually defined experiment name in tracking.py but didn't tell the same name in command so it created another experiment 3 conflicted with file name.
- Add "`--experiment-name Basic_Tracking-Demo`" at the end of line 2.
- Can use remote github repos (mlflow project) 3, parameterize dataset (Maybe I wanna use my own data)
- The "local" in line #2 tells mlflow to use current virtual environments. We can also set a default in mlflow project or give an environment (as configured in yaml) as parameter.

MLflow - Model.

- ⇒ Takes the trained model (with however many parameters) and from whomever framework (sklearn, Pytorch) and wraps it in a single wrapper (with the information needed by mlflow) to make mlflow model.
- ⇒ After it is saved, all of every model comes down to a single interface for inference & whatnot.

infer-signatures:

Models.py:

- 1) Made an sklearn model
 - 2) Made a Pytorch model
 - 3) Made a custom sklearn model with scaling. Same mlflow syntax.
-
- 4) Load sklearn model & predict
 - 5) " PyTorch (custom) model & predict.

MLFlow - Model Registry

• Provides centralized

(current)

1) Model store

2) Versioning

3) Life-cycle (None → Staging → Production → Archived)

Stage variable ↓

⇒ Files will now have logged models which appears on the dashboard but that's not the same as register (model-registry)

⇒ Registering model (for UI & managing repository).

1) Version

2) Stage

3) Tags

4) Aliases

⇒ Compare models

⇒ fetch & do whatever you want with all ~~data~~ models, their ~~farar~~ versions, parameters, metrics etc..

MLFlow Basic Run:

- 1) mlflow.set_experiment("Experiment_name")
- 2) with mlflow.start_run(run_name = "base_rf_model") as run:
 - 3) mlflow.log_params(params_dict)
 // create params dictionary
 - 4) mlflow.log_metrics(metrics_dict)
 // train model, predict, store metrics
 - 5) mlflow.set_tag("tag_type", tag_string)
 // create artifacts
 - 6) mlflow.log_artifact(Artifact_path)
 - 7) infer_signature(X_test, model.predict(X_test))
 - 8) mlflow.sklearn.log_model(model, "model", signature=signature, input_example=X_test[:5], metadata={
 "model_type": "RandomForest",
 "task": "binary-classification",
 "framework": "sklearn",
 "version": "1.0.0" })
- 1) Simple tracking
- 2) hyper-parameter tuning
- 3) batch logging
- 4) custom artifacts

ML Project

config.yaml

python-env.yaml

requirements.txt

train.py

ML Project:

name: mlflow-tutorial

python_env: python-env.yaml

entry_points:

main:

command: "python run-all.py"

tracking:

command: "python 1-tracking.py"

projects:

parameters:

n_estimators: {type: int, default: 10}

depth_max_depth: {type: int, default: 5}

command: "python 2-projects.py {n_estimators} {max_depth}"

models:

command: "python 3-models.py"

model_registry:

command: "python 4-model-registry.py"

PAPERWORK

python-env.yaml:

```
python: "3.8"  
build-dependencies:  
- "pip"  
dependencies: 2.8.0  
- mlflow>= 2.8.0  
- pandas>= 2.0.0
```

{

```
python: "3.8"  
build-dependencies:  
- "pip"  
dependencies:  
- mlflow >= 2.8.0
```

→ Once model is logged: run-id = run.info.run_id

```
model-uri = f"runid:{run_id}/model"
```

1) loaded_sklearn_model = mlflow.sklearn.load_model(model_uri)

2) loaded_pyfunc_model = mlflow.pyfunc.load_model(model_uri)

II Predictions

1) loaded_sklearn_model.predict(x-test[:5])

2) loaded_pyfunc_model.predict(x-test[:5])

3) loaded_pyfunc_model.predict(x-test-df)

Inheritance allows for custom preprocessing ↳ slice: mlflow.pyfunc.PythonModel

①

```
model_name = "classification-model"  
model_ifo = mlflow.sklearn.log_model(  
    model,  
    "model",  
    Signature=signature,  
    registered_model_name= model_name  
)
```

②

Just use this again if you want to edit, i.e ^{version} version 2.

③

If you keep the model name same you can even change the whole model, different algorithm and size. It'll just create a newer version.

④

```
client = MlflowClient()
```

```
client.search_model_versions({"name": "model-name"})
```

```
client.transition_model_version_stage(
```

```
    name="model-name",
```

```
    version=1,
```

```
    stage="Staging"
```

```
client.update_registered_model("adding description to model")
```

```
    name="model-name",
```

```
    description="....")
```

PAPERWORK

)

```
client.update_model_version( // updating version descriptions.  
    name = model_name,  
    version = 1,  
    description = "..."  
)
```

```
client.set_model_version_tag(model_name, 1, "algorithm", "RandomForest")  
    <variable> <value>
```

```
client.get_registered_model(model_name)
```

// loading models from registry: (by version)

```
model_version_uri = f"models:/ {model_name}/ 2"
```

```
loaded_model_v2 = mlflow.sklearn.load_model(model_version_uri)
```

```
predictions = loaded_model_v2.predict(x_test[:5])
```

// (by stage)

```
model_stage_uri = f"models:/ {model_name}/ Production"
```

// (latest_version)

Compare models 3 aliasing!

APACHE AIRFLOW

11-10-2022

A tool for authoring, scheduling, and monitoring

2) Directed Acyclic Graphs (DAG). Workflows. Airflow Task

=> Task: Instantiation of an operator.

=> Operator / Sensor / Hook : Operator performs actions (Bash Operator, Python Operator).

Sensors wait for external conditions.

Hooks implement connections to external service.

=> Scheduler: Schedules DAG runs & creates task instances.

=> Executor: Controls parallelism and how tasks run (Sequential, Local Celery, K8s etc.)

=> XCom: Light-weight per-task-instance communication system

=> Variables & Connections: Global config & credentials storage.

APACHE AIRFLOW

⇒ Better way to run is Docker-Compose.

Docker-Compose:

- 1) Postgres (for db)
- 2) x-airflow-common (doesn't get executed by docker-compose, to use somewhere else in the file)
- 3) airflow-scheduler
- 4) airflow-init

Logs Folder:

Logs FOLDER:

⇒ Everything about logs is logged over here

Demo: (UI)

1) Added DAG file manually to dag folder (vs code)
will reflect on a airflow UI.

2) DAG runs. Decoding the UI.

3) Run vs Parse DAG

1) x Communication.

DAG Code: (The file added manually to dag folder)

1) Imports

2) default_args.

3) with DAG (as

dag_id:

default_args: default_args

...

) as dag:

start_task = BashOperator (

task_id = "start"

bash_command = "'echo \"Starting pipeline at \$(date)\"'"
,

process_task = BashOperator (

...

)

complete_task = BashOperator (

...

)

start_task >> process_task >> complete_task (dependencies linear)

PAPERWORK

Example - 2 : Code file (DAG) (ETL example to show dependency)

- 1) Uses decorators `@dag` & `@task`.
- 2) Define tasks as functions using `@task`.
- 3) Define dependencies implicitly by just using tasks as functions and feeding the returned data forward.

(i)

→ → →

- a) Seeing it on UI
- b) x Con display on UI
- c) Showing & explaining automatic & manual methods of x Con (on Readme).

AIRFLOW

Task Dependencies:

- Linear

task1 \Rightarrow task2 \Rightarrow task3

1
2
1
3

- Multiple upstream tasks

task1 \Rightarrow task3

task2 \Rightarrow task3

1 2
 \ /
 3

- Multiple downstream tasks:

task1 \Rightarrow [task2, task3, task4]

2 1 \ /
3 4

- Complex dependencies.

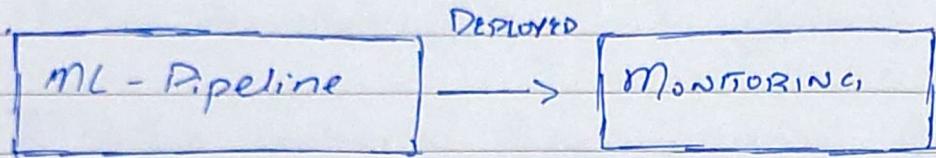
$\xrightarrow{\quad}$ $\xrightarrow{\quad}$ $\xrightarrow{\quad}$

\Rightarrow Connections

\Rightarrow Holes

PROMETHEUS & GRAPHANA - MONITORING

20-11-2025



Prometheus: Getting data from the application / micro-service cluster

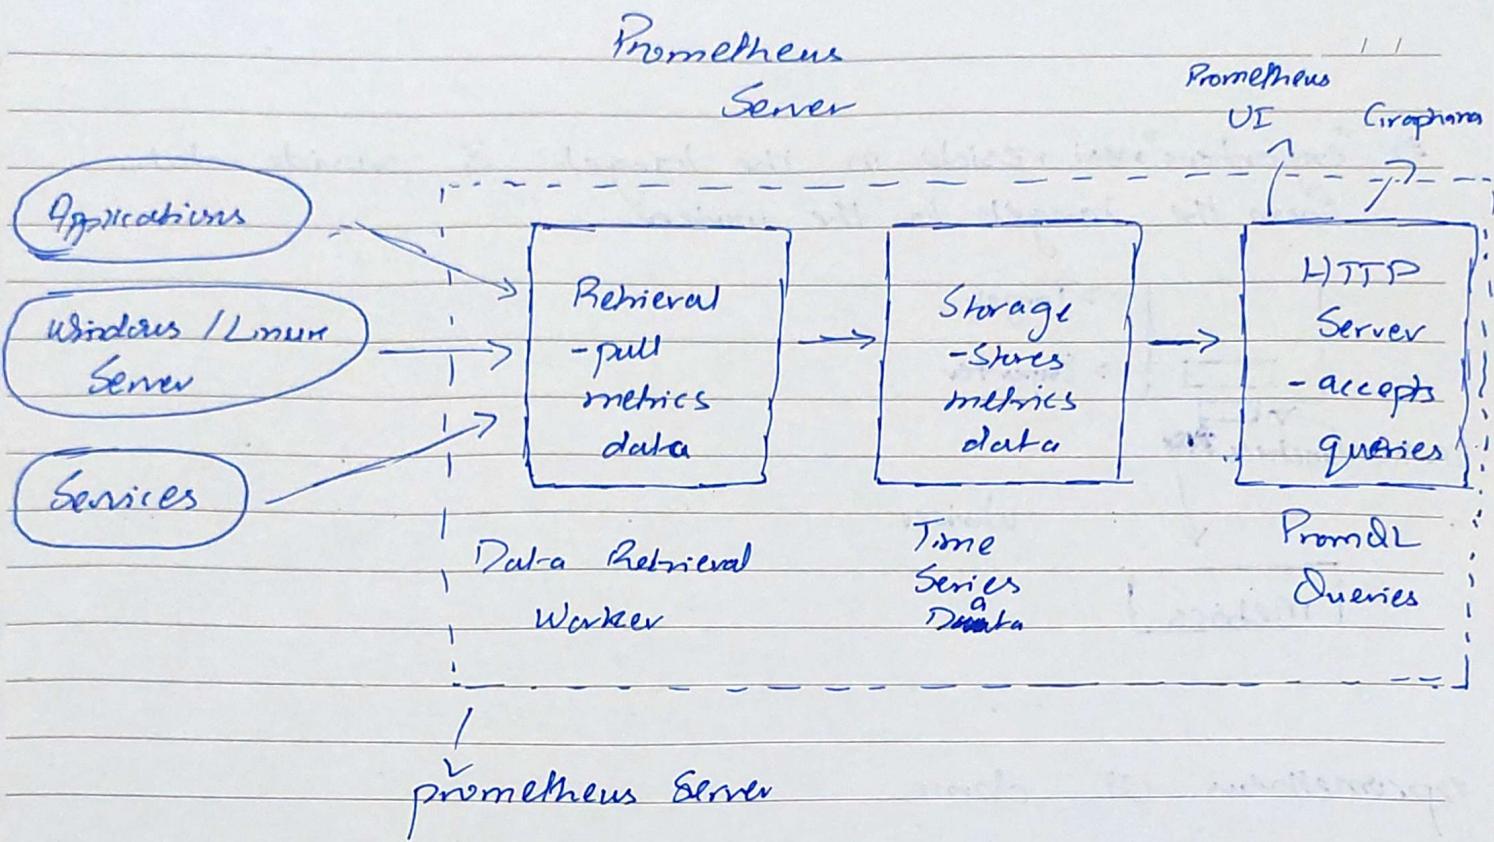
Graphana:

Prometheus:

- ⇒ Pull based data scraping
- ⇒ Multi-dimensional data labels
- ⇒ Use Prom QL (The query language used by Prometheus to query data)
- ⇒ Local Timeseries storage
- ⇒ Integration with alert manager with alerts.

Graphana:

- ⇒ Interactive Dashboard
 - ⇒ Time series Visualization
 - ⇒ Alerts
 - ⇒ Integrates with multiple data sources (including prometheus.)
- ⇒ We use docker-compose to run these together (prometheus, graphana, alert-manager)

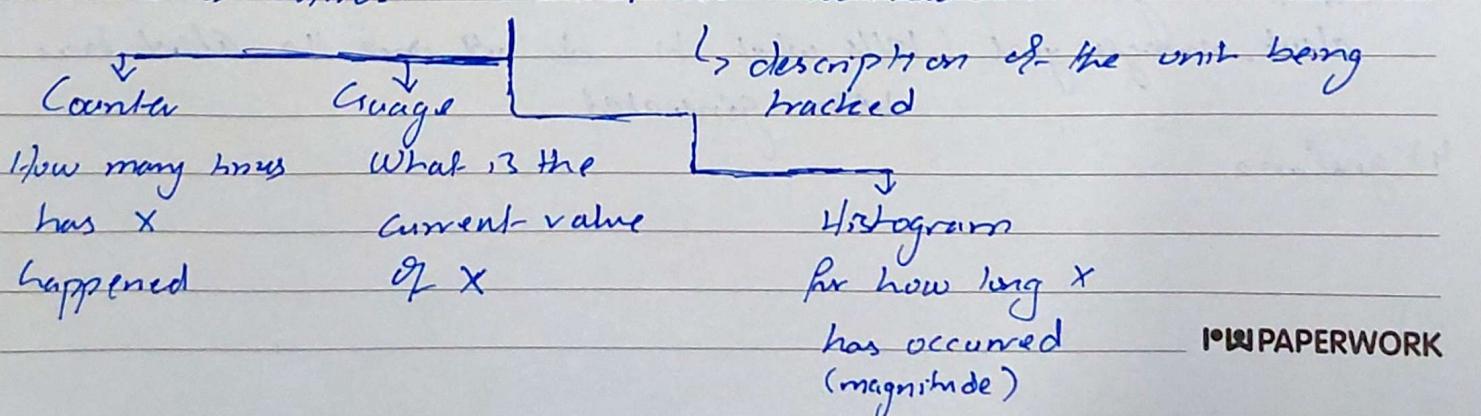


Targets:

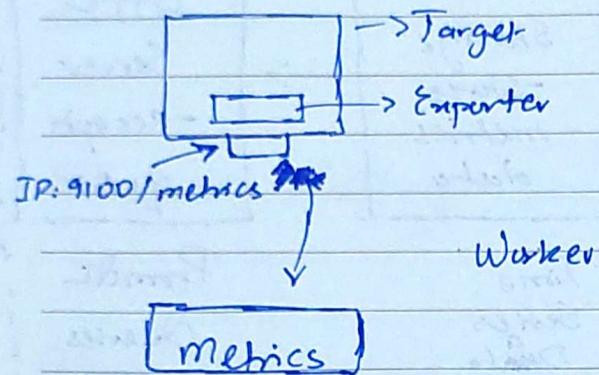
- ⇒ What Sources does prometheus monitor
- 1) Linux / Windows Server
- 2) Applications
- 3) Services

Metrics:

- ⇒ What units are monitored from those targets.
- ⇒ Format: Human readable text
- ⇒ Metrics Entries: TYPE & HELP attributes.



⇒ "exporters" reside in the target. To provide data from the target to the worker.



⇒ prometheus UI demo

25 - 11 - 2025

Prometheus Code Demo:

1) docker-compose.yml

2) prometheus:

prometheus.yml (all the endpoints running / exporters etc.)
alert.rules.yml (conditions to generate alerts)

3) alert_manager

alert_manager.yml (tells what to do once the alert has been generated)

4) grafana

5)

Date

PromQL

1) variable_name (execute)

2) variable_name {idle mode = 'idle'} → Filter by parameter value

3)

CRAFPHANA

- 1) Install prometheus plugin.
- 2) Add new connection

unlimited res.
CPU usage max.

/dev/null

yes > /dev/null

Date

alert-rules.yaml

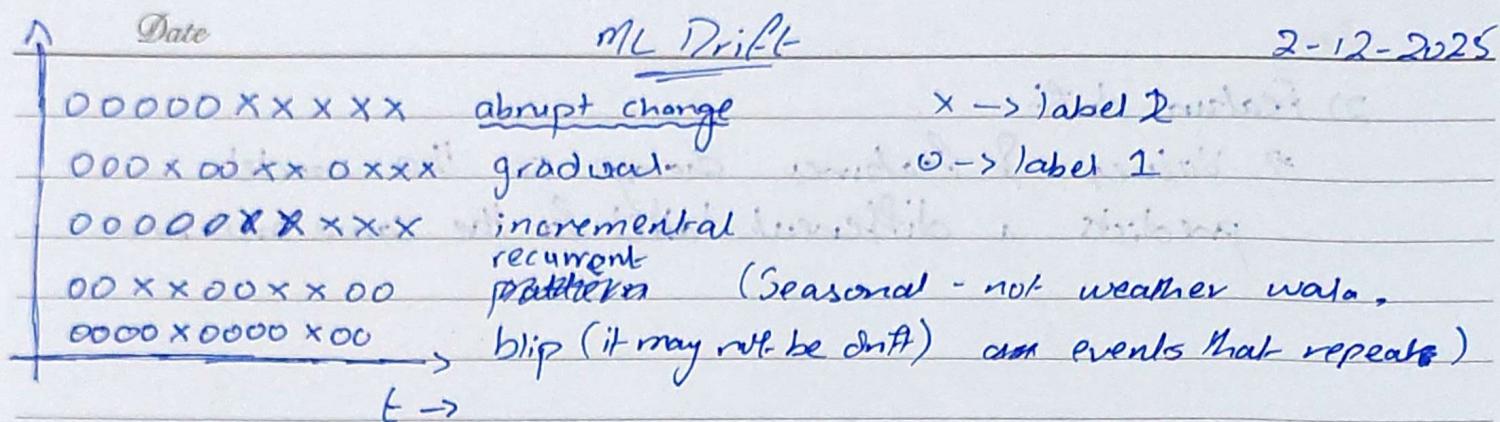
- => Gives the PromQL query to run
- => The ~~for~~ time that will be waited before generating alerts
(to remove false positives)
- => Gives summary & description of alerts.
- => Blackbox: An alert definition that checks if endpoints are up.

alert-manager demo

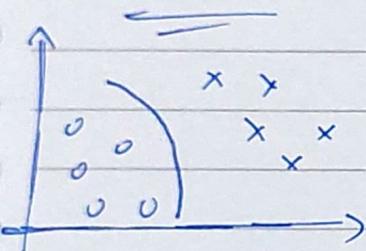
connection with slack

details on github (will be uploaded by sru)

Prediction



Concept Drift:



→ For e.g. a model learnt this boundary based on the data / label values.

⇒ Now we change a concept / condition in the real world that changes the labels on those same data.

⇒ So our data didn't change but the label formation did and hence we need a new ^{boundary} line (retraining)

⇒ We change a condition that people below 1800 work instead of 1500 will not be given a loan

Data Drift:

1) LABEL DRIFT

2) FEATURE DRIFT

1) LABEL DRIFT:

⇒ We didn't change any concept

⇒ The values of the input didn't change

⇒ We noticed an increase in the occurrences of labels. Maybe because the instances of inputs increased.

⇒ More people applying for loans, (More rejected & accepted). Labels for both will drift

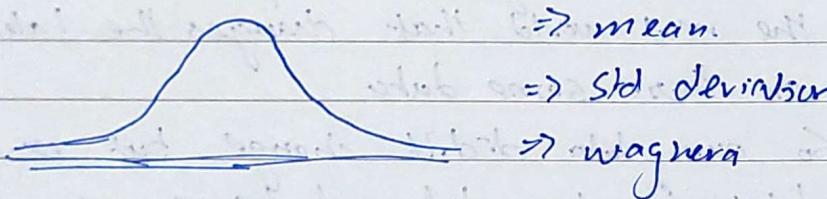
Date

2) Feature drift.

→ Values of features change so the model predicts a different label for the same instance.

Detecting Drift In Unsupervised Learning

1) Numerical Variables



2) Categorical

→ Frequency Trends



3) Relations b/w numerical variables

⇒ pearson (r^2)

4) Relation b/w

F_1 & F_2 two Features can have low correlation (r^2) but still be strongly linked (quadratically or any other non-linear relation) since r^2 is just a linear correlation mapping.